

**Note to other teachers and users of these slides:** We would be delighted if you found our material useful for giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. If you make use of a significant portion of these slides in your own lecture, please include this message, or a link to our web site: <http://www.mmnds.org>

# CS246: Mining Massive Data Sets

## Intro, MapReduce & Spark

CS246: Mining Massive Data Sets  
Jure Leskovec, Stanford University  
<http://cs246.stanford.edu>





Data contains value and knowledge

# Data Mining

- But to extract the knowledge data needs to be
  - Stored (systems)
  - Managed (databases)
  - And **ANALYZED** ← this class

**Data Mining  $\approx$  Big Data  $\approx$   
Predictive Analytics  $\approx$   
Data Science  $\approx$  Machine Learning**

# What This Course Is About

- ***Data mining*** = extraction of actionable information from (usually) very large datasets, is the subject of extreme hype, fear, and interest
- It's not all about machine learning
- But most of it is
- Emphasis in CS246 on algorithms that **scale**
  - Parallelization often essential

# Data Mining Methods

- **Descriptive methods**


- Find human-interpretable patterns that describe the data
  - **Example:** Clustering

- **Predictive methods**

- Use some variables to predict unknown or future values of other variables
  - **Example:** Recommender systems

# This Class: CS246

- This combines best of machine learning, statistics, artificial intelligence, databases but more stress on
  - **Scalability** (big data)
  - **Algorithms**
  - **Computing architectures**
  - Automation for handling **large data**



# What will we learn?

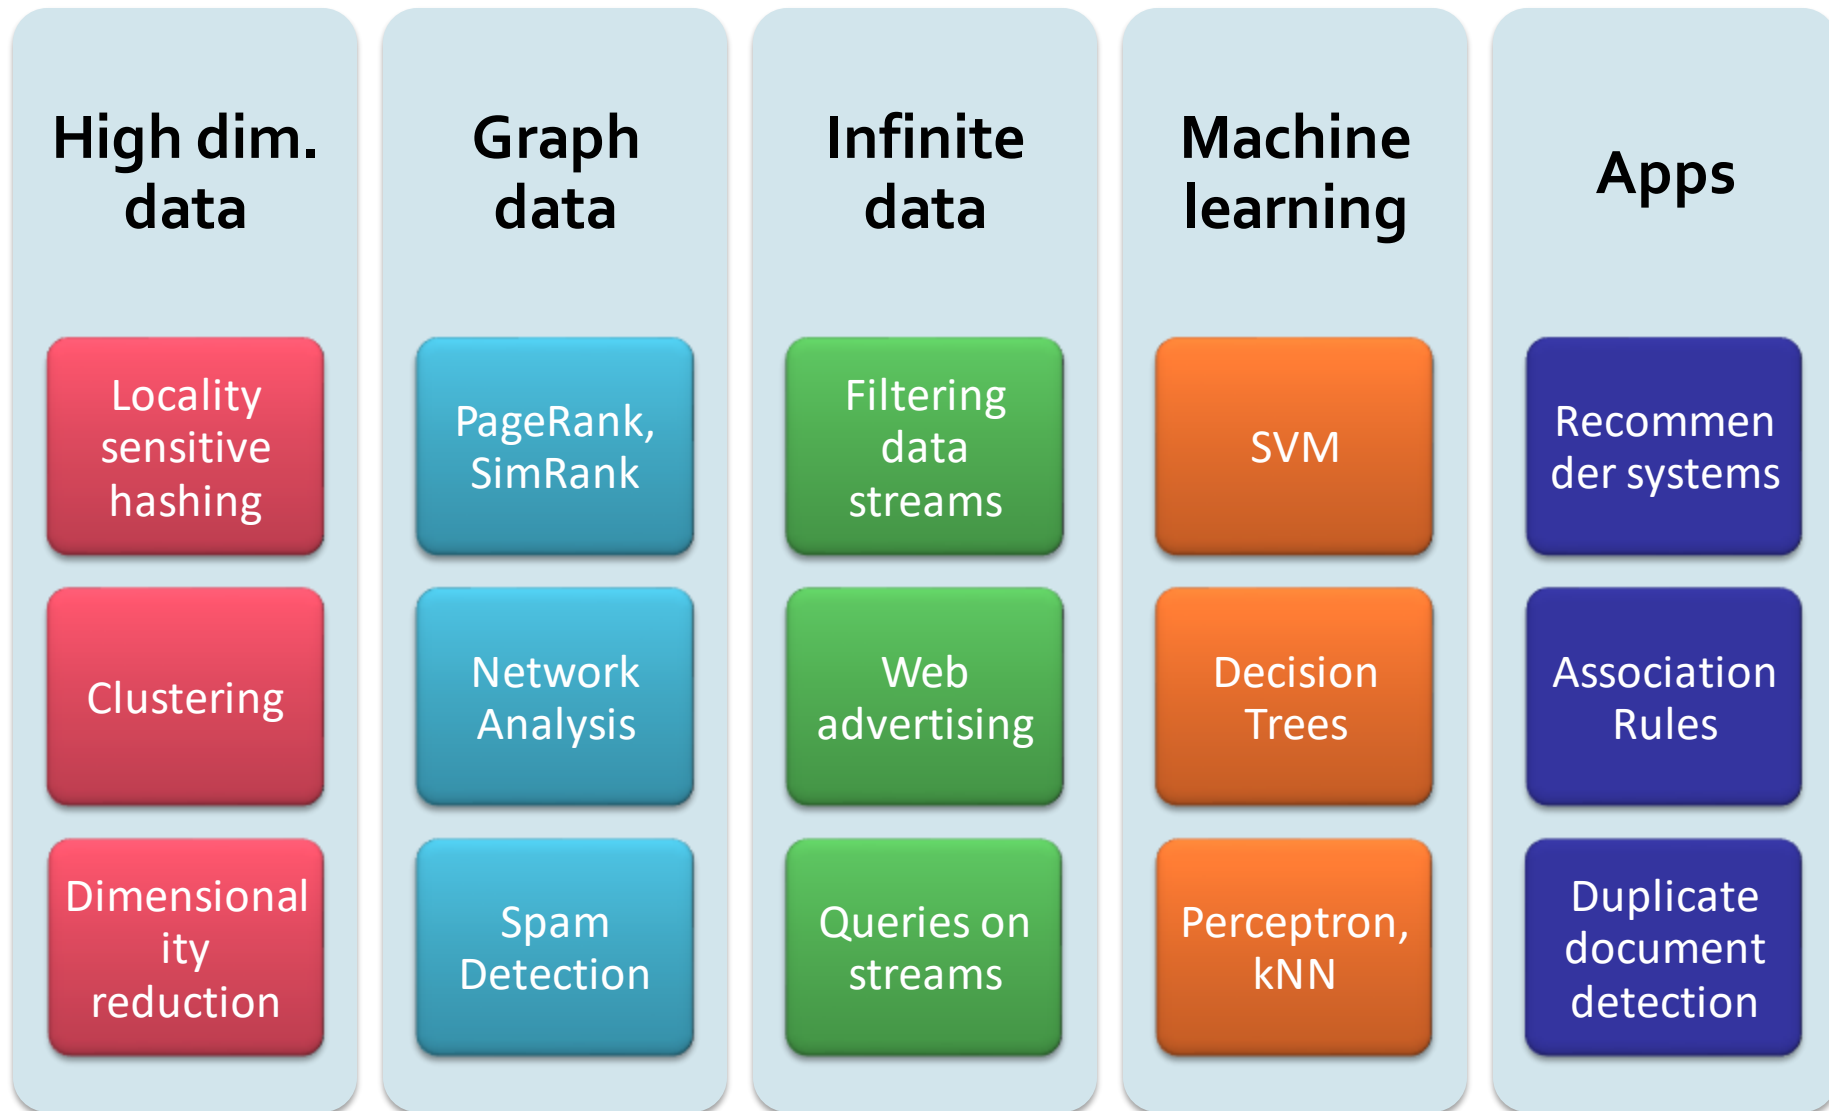
- **We will learn to mine different types of data:**
  - Data is high dimensional
  - Data is a graph
  - Data is infinite/never-ending
  - Data is labeled
- **We will learn to use different models of computation:**
  - MapReduce
  - Streams and online algorithms
  - Single machine in-memory

# What will we learn?

- **We will learn to solve real-world problems:**
  - Recommender systems
  - Market Basket Analysis
  - Spam detection
  - Duplicate document detection
- **We will learn various “tools”:**
  - Linear algebra (SVD, Rec. Sys., Communities)
  - Optimization (stochastic gradient descent)
  - Dynamic programming (frequent itemsets)
  - Hashing (LSH, Bloom filters)



# How the Class Fits Together





# How do you want that data?

# Course Logistics

# Course Staff

## Instructor



Jure Leskovec

## Course Coordinator



Natasha Sharp

## Teaching Assistants



Alexandra (Alex) Porter  
Head TA



Jiaxuan You



Yanjun Chen



Gordon (Trey) Connelly



Tsun-Han (Jerry)  
Huang

# Course Format: Zoom&Canvas

**Lectures: Tue/Thu 12:30-1:50pm PST**

Live on Zoom, recording available on Canvas

- **60 min lecture:**

- If you have a clarification question, post it in Zoom chat, TAs will answer

- **20 min Q&A:**

- Post a question in Zoom chat, Jure will answer and discuss

# Logistics: Communication

- **Piazza:**

- <http://piazza.com/stanford/spring2021/cs246>

- **Use Piazza for all questions and public communication**

- Search the feed before asking a duplicate question
- Please tag your posts and please no one-liners

- **For e-mailing course staff always use:**

- [cs246-spr2021-staff@lists.stanford.edu](mailto:cs246-spr2021-staff@lists.stanford.edu)

- **We will post course announcements to Piazza (hence check it regularly!)**

**Auditors are welcome!**

(please send request to [nsharp@stanford.edu](mailto:nsharp@stanford.edu) to add you to Canvas)

# Resources

- **Course website:** <http://cs246.stanford.edu>
  - Lecture slides (at least 30min before the lecture)
  - Homework, solutions, readings posted on Piazza
- **Class textbook:** **Mining of Massive Datasets** by A. Rajaraman, J. Ullman, and J. Leskovec
  - Sold by Cambridge Uni. Press but available for free at <http://mmds.org>
- **MOOC:** [www.youtube.com/channel/UC\\_Oao2FYkLAUIUVkBfze4jg/videos](http://www.youtube.com/channel/UC_Oao2FYkLAUIUVkBfze4jg/videos)

# CS246 Office Hours

- **Office hours:**

- See course website <http://cs246.stanford.edu> for TA office hours
  - **We start Office Hours this Friday!**
- Office hours will be held on Zoom and use [QueueStatus](#)
  - Links will be posted on Piazza and Canvas



# Recitation Sessions

- Videos and materials on Canvas
- **Spark tutorial:**
  - Video
  - Follows Colab 0
- **Review of basic probability and proof techniques**
  - Video and handout
- **Review of linear algebra**
  - Video and handout

# Work for the Course: Homework

- **4 longer homeworks: 60%**
  - Four major assignments, involving programming, proofs, algorithm development.
  - Assignments take lots of time (+20h). **Start early!!**
- **How to submit?**
  - **Homework write-up:**
    - Submit via [Gradescope](#)
    - Enroll to CS246 on Canvas, and you will be automatically added to the course Gradescope
  - **Homework code:**
    - If the homework requires a code submission, you will find a separate assignment for it on Gradescope, e.g., **HW1 (Code)**
    - **Forgetting to submit code will result in point deduction.**

# Homework Calendar

## ■ Homework schedule:

Date (23:59 PST)	Out	In
04/01, Thu	HW1	
04/15, Thu	HW2	HW1
04/29, Thu	HW3	HW2
05/13, Thu	HW4	HW3
05/27, Thu		HW4

## ■ Two late periods for HWs for the quarter:

- Late period expires on the following Monday 23:59 PST
- Can use max 1 late period per HW

# Work for the Course: Colabs

- **Short weekly Colab notebooks: 40%**
  - Colab notebooks are posted every **Thursday**
    - 10 in total, from 0 to 9, each worth 4%
  - Due one week later on **Thursday 23:59 PST. No late days!**
    - First 2 Colabs will be posted on Thu, including detailed submission instructions to Gradescope
    - Colab 0 (Spark Tutorial) is solved step-by-step in the [Spark Recitation video](#).
  - Colabs require around **1hr of work**.
    - **And a few lines of code.**
  - “Colab” is a **free cloud service from Google**, hosting [Jupyter notebooks](#) with free access to GPU and TPU

# Work for the Course: Final Exam

- **NO Final exam**
- **Extra credit:** Proportional to your contribution (up to 2%)
  - Course attendance, asking questions, discussion
  - For participating in Piazza discussions
    - Especially valuable are answers to questions posed by other students
  - Reporting bugs in course materials

# Prerequisites

- **Programming:** Python or Java
- **Basic Algorithms:** CS161 is surely sufficient
- **Probability:** e.g., CS109 or Stats116
  - There will be a review session and a review doc is linked from the class home page
- **Linear algebra:**
  - Another review doc + review session is available
- **Multivariable calculus**
- **Database systems** (SQL, relational algebra):
  - CS145 is sufficient but not necessary

# What If I Don't Know All This Stuff?

- **Each of the topics listed is important for a small part of the course:**
  - If you are missing an item of background, you could consider just-in-time learning of the needed material
- **The exception is programming:**
  - To do well in this course, you really need to be comfortable with writing code in Python or Java

# Honor Code

- **We'll follow the standard CS Dept. approach:**  
You can get help, but you ***MUST*** acknowledge the help on the work you hand in
- Failure to acknowledge your sources is a ***violation of the Honor Code***
- We use MOSS to check the originality of your code



# Honor Code – (2)

- You can talk to others about the algorithm(s) to be used to solve a homework problem;
  - As long as you then mention their name(s) on the work you submit.
- You should not use code of others or be looking at code of others when you write your own:
  - (don't search/post code on Github, and similar)
  - You can talk to people but have to write your own solution/code
  - If you fail to mention your sources, MOSS will catch it, which will result in an HC violation.

# What's After the Class

- **CS341: Project in Data Mining (2021/22)**
  - Research project on Big Data
  - Groups of 3 students
  - We provide interesting data, computing resources (Google Cloud) and mentoring
- **My group has RA positions open:**
  - See <http://snap.stanford.edu/apply/>
- **In past years we used to run CS246H.  
We won't be able to run CS246H this year.**

# Final Thoughts

- **CS246 is fast paced!**
  - Requires programming maturity
  - Strong math skills
    - SCPD students tend to be rusty on math/theory
- **Course time commitment:**
  - Homeworks take +20h
  - Colab notebooks take about 1h
- Form study groups
- **It's going to be fun and hard work. 😊**

# Distributed Computing for Data Mining



# Large-scale Computing

- Large-scale computing for data mining problems on commodity hardware
- Challenges:
  - How do you distribute computation?
  - How can we make it easy to write distributed programs?
  - Machines fail:
    - One server may stay up 3 years (1,000 days)
    - If you have 1,000 servers, expect to lose 1/day
    - With 1M machines 1,000 machines fail every day!

# An Idea and a Solution

- **Issue:**

**Copying data over a network takes time**

- **Idea:**

- Bring computation to data
- Store files multiple times for reliability

- **Spark/Hadoop** address these problems

- **Storage Infrastructure – File system**

- Google: GFS. Hadoop: HDFS

- **Programming model**

- MapReduce
- Spark

# Storage Infrastructure

- **Problem:**

- If nodes fail, how to store data persistently?

- **Answer:**

- **Distributed File System**

- Provides global file namespace

- **Typical usage pattern:**

- Huge files (100s of GB to TB)
  - Data is rarely updated in place
  - Reads and appends are common



# Distributed File System

## ■ **Chunk servers**

- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

## ■ **Master node**


- a.k.a. Name Node in Hadoop's HDFS
- Stores metadata about where files are stored
- Might be replicated

## ■ **Client library for file access**

- Talks to master to find chunk servers
- Connects directly to chunk servers to access data

# Distributed File System

- **Reliable distributed file system**
- Data kept in “chunks” spread across machines
- Each chunk **replicated** on different machines
  - Seamless recovery from disk or machine failure



Bring computation directly to the data!

Chunk servers also serve as compute servers

# MapReduce: Early Distributed Computing Programming Model

# Programming Model: MapReduce

- MapReduce is a **style of programming** designed for:
  1. Easy parallel programming
  2. Invisible management of hardware and software failures
  3. Easy management of very-large-scale data
- It has several **implementations**, including Hadoop, Spark (used in this class), Flink, and the original Google implementation just called “MapReduce”

# MapReduce: Overview

## 3 steps of MapReduce

### ■ Map:

- Apply a user-written *Map function* to each input element
  - *Mapper* applies the Map function to a single element
    - Many mappers grouped in a *Map task* (the unit of parallelism)
- The output of the Map function is a set of 0, 1, or more *key-value pairs*.

### ■ Group by key: Sort and shuffle

- System sorts all the key-value pairs by key, and outputs key-(list of values) pairs

### ■ Reduce:

- User-written *Reduce function* is applied to each key-(list of values)


Outline stays the same, **Map** and **Reduce** change to fit the problem

# Map-Reduce: A diagram

## MAP:

Read input and produces a set of key-value pairs


Input




Intermediate

## Group by key:

Collect all pairs with same key  
(Hash merge, Shuffle, Sort, Partition)



Grouped

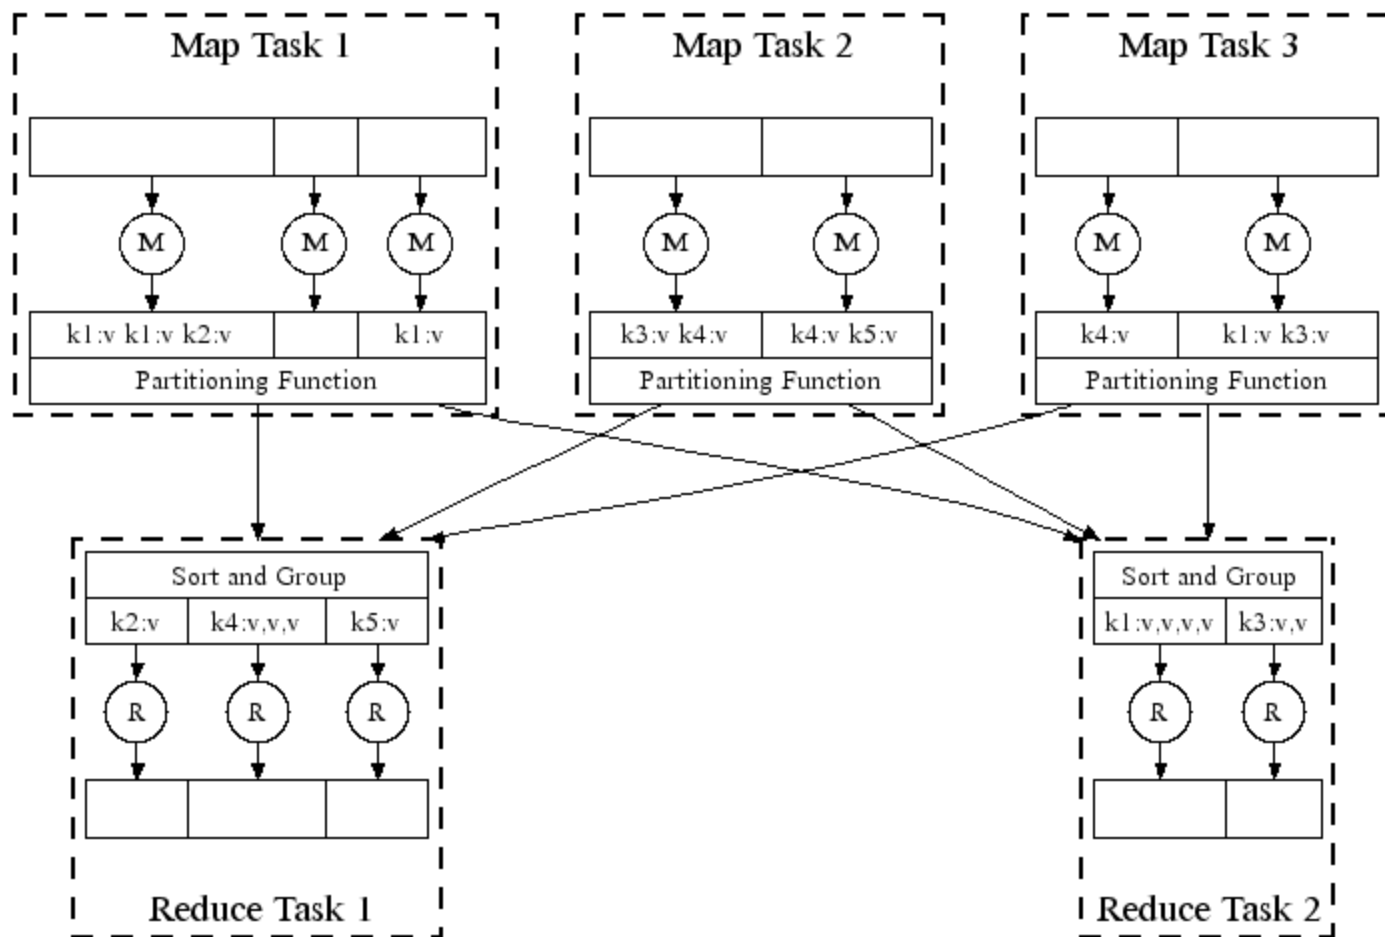


Output

## Reduce:


Collect all values belonging to the key and output

# Map-Reduce: In Parallel



All phases are distributed with many tasks doing the work

# MapReduce Pattern





# Example: Word Counting

## Example MapReduce task:

- We have a huge text document
- Count the number of times each distinct word appears in the file
- **Many applications of this:**
  - Analyze web server logs to find popular URLs
  - Statistical machine translation:
    - Need to count number of times every 5-word sequence occurs in a large corpus of documents

# MapReduce: Word Counting

Provided by the  
programmer

## MAP:

Read input and  
produces a set of  
key-value pairs

(The, 1)

(crew, 1)

(of, 1)

(the, 1)

(space, 1)

(shuttle, 1)

(Endeavor, 1)

(recently, 1)

....

(key, value)

## Group by key:

Collect all pairs  
with same key

(crew, 1)

(crew, 1)

(space, 1)

(the, 1)

(the, 1)

(the, 1)

(shuttle, 1)

(recently, 1)

...

(key, value)

Provided by the  
programmer

## Reduce:

Collect all values  
belonging to the  
key and output

(crew, 2)

(space, 1)

(the, 3)

(shuttle, 1)

(recently, 1)

...

(key, value)

Only sequential reads

The crew of the space shuttle Endeavor recently returned to Earth as ambassadors, harbingers of a new era of space exploration. Scientists at NASA are saying that the recent assembly of the Dextre bot is the first step in a long-term space-based man/machine partnership. "The work we're doing now -- the robotics we're doing - is what we're going to need .....

Big document

# Word Count Using MapReduce

**map(key, value):**

```
# key: document name; value: text of the document
  for each word w in value:
    emit(w, 1)
```

**reduce(key, values):**

```
# key: a word; value: an iterator over counts
  result = 0
  for each count v in values:
    result += v
  emit(key, result)
```

# MapReduce: Environment

## MapReduce environment takes care of:

- **Partitioning** the input data
- **Scheduling** the program's execution across a set of machines
- Performing the **group by key** step
  - In practice this is the bottleneck
- Handling machine **failures**
- Managing required inter-machine **communication**

# Dealing with Failures

## ■ Map worker failure


- Map tasks completed or in-progress at worker are reset to idle and rescheduled
- Reduce workers are notified when map task is rescheduled on another worker

## ■ Reduce worker failure

- Only in-progress tasks are reset to idle and the reduce task is restarted

# Spark: Extends MapReduce

# Problems with MapReduce



## ■ MapReduce:

- Incurs substantial overheads due to data replication, disk I/O, and serialization

# Problems with MapReduce

- **Two major limitations of MapReduce:**
  - Difficulty of programming directly in MapReduce
    - Many problems aren't easily described as map-reduce
  - Performance bottlenecks, or batch not fitting the use cases
    - Persistence to disk typically slower than in-memory work
- **In short, MapReduce doesn't compose well for large applications**
  - Many times, one needs to chain multiple map-reduce steps.



# Data-Flow Systems

- **MapReduce uses two “ranks” of tasks:**  
One for **Map** the second for **Reduce**
  - Data flows from the first rank to the second
- **Data-Flow Systems generalize this in two ways:**
  1. Allow any number of tasks/ranks
  2. Allow functions other than Map and Reduce
  - As long as data flow is in one direction only, we can have the blocking property and allow recovery of tasks rather than whole jobs

# Spark: Most Popular Data-Flow System

- **Expressive computing system, not limited to the map-reduce model**
- **Additions to MapReduce model:**
  - Fast data sharing
    - Avoids saving intermediate results to disk
    - Caches data for repetitive queries (e.g. for machine learning)
  - General execution graphs (DAGs)
  - Richer functions than just map and reduce
- **Compatible with Hadoop**

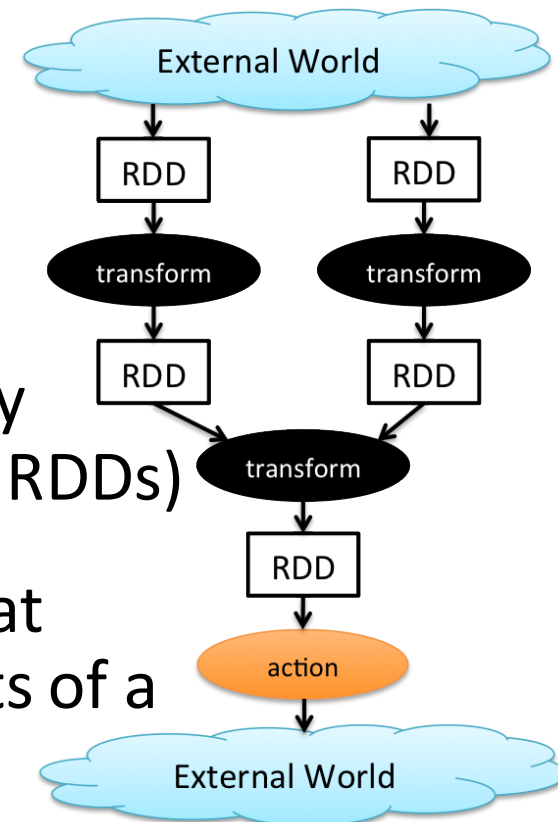
# Spark: Overview

- Open source software (Apache Foundation)
- Supports **Java, Scala and Python**
- **Key construct/idea:** Resilient Distributed Dataset (RDD)
- **Higher-level APIs:** DataFrames & DataSets
  - Introduced in more recent versions of Spark
  - Different APIs for aggregate data, which allowed to introduce SQL support

# Spark: RDD

## Key concept *Resilient Distributed Dataset* (RDD)


- Partitioned collection of records
  - Generalizes (key-value) pairs
- Spread across the cluster, Read-only
- Caching dataset in memory
  - Different storage levels available
  - Fallback to disk possible
- RDDs can be created from Hadoop, or by transforming other RDDs (you can stack RDDs)
- RDDs are best suited for applications that apply the same operation to all elements of a dataset



# Spark RDD Operations

- **Transformations** build RDDs through deterministic operations on other RDDs:
  - Transformations include *map, filter, join, union, intersection, distinct*
  - **Lazy evaluation:** Nothing computed until an action requires it
- **Actions** to return value or export data
  - Actions include *count, collect, reduce, save*
  - Actions can be applied to RDDs; actions force calculations and return values

# Task Scheduler: General DAGs



- Supports general task graphs
- Pipelines functions where possible
- Cache-aware data reuse & locality
- Partitioning-aware to avoid shuffles

# DataFrame & Dataset

## ■ DataFrame:

- Unlike an RDD, data organized into named columns, e.g. a **table in a relational database**.
- Imposes a structure onto a distributed collection of data, allowing higher-level abstraction

## ■ Dataset:

- Extension of DataFrame API which provides **type-safe, object-oriented programming interface** (compile-time error detection)

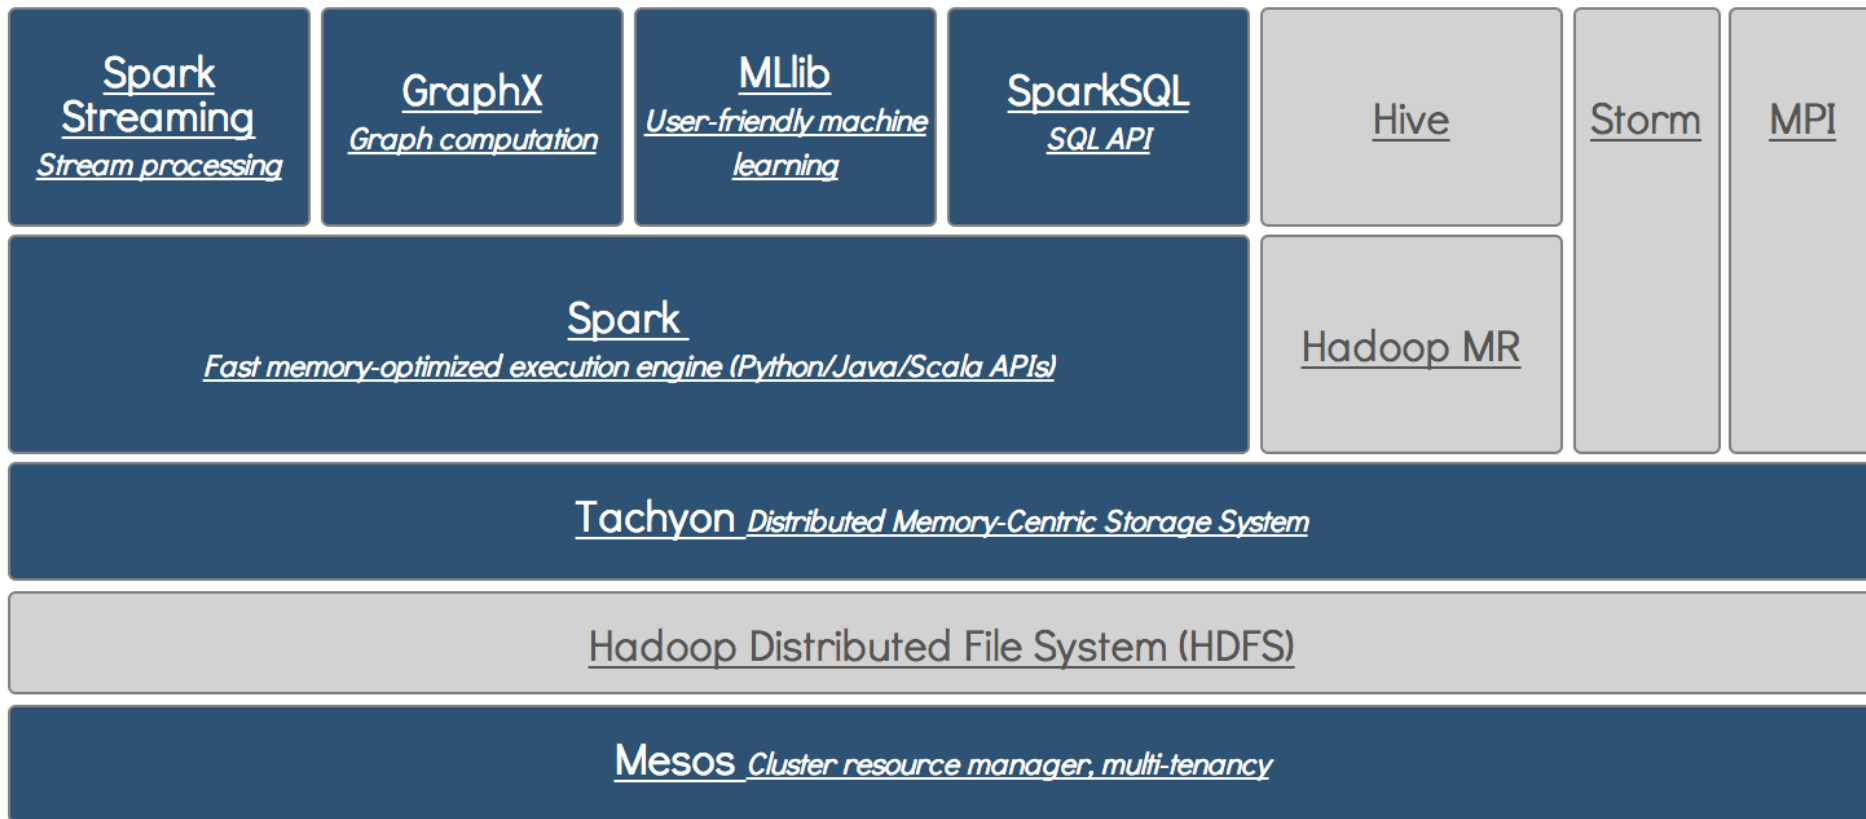
Both built on Spark SQL engine. Both can be converted back to an RDD.

# Useful Libraries for Spark

- Spark SQL
- Spark Streaming – **stream processing of live datastreams**
- MLlib – **scalable machine learning**
- GraphX – **graph manipulation**
  - Extends Spark RDD with Graph abstraction: a directed multigraph with properties attached to each vertex and edge



# Data Analytics Software Stack



# Spark vs. Hadoop MapReduce

- Performance: **Spark normally faster** but **with caveats**
  - Spark can process data in-memory; Hadoop MapReduce persists back to the disk after a map or reduce action
  - Spark generally outperforms MapReduce, but it **often needs lots of memory to perform well**; if there are other resource-demanding services or can't fit in memory, Spark degrades
  - MapReduce easily runs alongside other services with minor performance differences, & works well with the 1-pass jobs it was designed for
- Ease of use: **Spark is easier to program** (higher-level APIs)
- Data processing: **Spark more general**

# Problems Suited for MapReduce

# Example: Host size

- **Suppose we have a large web corpus**
- Look at the metadata file
  - Lines of the form: (URL, size, date, ...)
- **For each host, find the total number of bytes**
  - That is, the sum of the page sizes for all URLs from that particular host
- **Other examples:**
  - Link analysis and graph processing
  - Machine Learning algorithms

# Example: Language Model

- **Statistical machine translation:**
  - Need to count number of times every 5-word sequence occurs in a large corpus of documents
- **Very easy with MapReduce:**
  - **Map:**
    - Extract (5-word sequence, count) from document
  - **Reduce:**
    - Combine the counts

# Example: Join By Map-Reduce

- Compute the natural join  $R(A,B) \bowtie S(B,C)$
- $R$  and  $S$  are each stored in files
- Tuples are pairs  $(a,b)$  or  $(b,c)$

A	B
$a_1$	$b_1$
$a_2$	$b_1$
$a_3$	$b_2$
$a_4$	$b_3$

R



B	C
$b_2$	$c_1$
$b_2$	$c_2$
$b_3$	$c_3$

S



A	C
$a_3$	$c_1$
$a_3$	$c_2$
$a_4$	$c_3$

# Map-Reduce Join

- Use a hash function  $h$  from B-values to  $1...k$
- **A Map process turns:**
  - Each input tuple  $R(a,b)$  into key-value pair  $(b,(a,R))$
  - Each input tuple  $S(b,c)$  into  $(b,(c,S))$
- **Map processes** send each key-value pair with key  $b$  to Reduce process  $h(b)$ 
  - Hadoop does this automatically; just tell it what  $k$  is.
- Each **Reduce process** matches all the pairs  $(b,(a,R))$  with all  $(b,(c,S))$  and outputs  $(a,b,c)$ .

# Problems NOT suitable for MapReduce

- **MapReduce is great for:**
  - Problems that require sequential data access
  - Large batch jobs (**not** interactive, real-time)
- **MapReduce is inefficient for problems where random (or irregular) access to data required:**
  - Graphs
  - Interdependent data
    - Machine learning
    - Comparisons of many pairs of items



# Cost Measures for Algorithms

- In MapReduce we quantify the cost of an algorithm using
  1. *Communication cost* = total I/O of all processes
  2. *Elapsed communication cost* = max of I/O along any path
  3. (*Elapsed*) *computation cost* analogous, but count only running time of processes

Note that here the big-O notation is not the most useful (adding more machines is always an option)

# Example: Cost Measures

- **For a map-reduce algorithm:**
  - **Communication cost** = input file size +  $2 \times$  (sum of the sizes of all files passed from Map processes to Reduce processes) + the sum of the output sizes of the Reduce processes.
  - **Elapsed communication cost** is the sum of the largest input + output for any map process, plus the same for any reduce process

# What Cost Measures Mean

- Either the I/O (communication) or processing (computation) cost dominates
  - Ignore one or the other
- Total cost tells what you pay in rent from your friendly neighborhood cloud
- Elapsed cost is wall-clock time using parallelism

# Cost of Map-Reduce Join

- **Total communication cost**  
 $= O(|R| + |S| + |R \bowtie S|)$
- **Elapsed communication cost**  $= O(s)$ 
  - We're going to pick  $k$  and the number of Map processes so that the I/O limit  $s$  is respected
  - We put a limit  $s$  on the amount of input or output that any one process can have.  **$s$  could be:**
    - What fits in main memory
    - What fits on local disk
- With proper indexes, computation cost is linear in the input + output size
  - So, computation cost is like communication cost